



AI Scribe Evidence File

Documentation evidence is not a cash-return finding.

Companion to "89.3% Semantic Agreement Is Not ROI: The Hospital AI Scribe Scorecard." This file separates the reported evidence from the proposed local investment decision.

What the accessible evidence reports

The September 23, 2026 Scientific Reports retrospective study concerns a Spanish outpatient network: more than 2.3 million consultations across 45 specialties. Detailed quality assessment covers 300 reports across 15 specialties, not the whole consultation population.

The publisher-indexed abstract reports 89.3% weighted mean semantic agreement and better readability and structured quality. It reports no measured labor savings, net financial return or patient outcomes. That absence is established only for the abstract, not for the entire paper. These are associations, not established causation.

The indexed competing-interest declaration states that all authors work for the network developing and evaluating the system and declare no personal financial interests in the technology. This is a disclosed relationship, not a finding that the results are invalid.

Direct primary-source routes

Publisher article

<https://www.nature.com/articles/s41598-026-72180-z>

DOI

<https://doi.org/10.1038/s41598-026-72180-z>

Publisher PDF route (access not established)

<https://www.nature.com/articles/s41598-026-72180-z.pdf>

Verification recorded September 27, 2026: publisher-indexed abstract, metadata and disclosure were accessible in the approved source record. Direct HTML/PDF access failed. Full-text methods and limitations were not verified. Links are primary-source routes, not an assertion of successful retrieval.



What This Evidence Cannot Decide

01 Similarity is not clinical accuracy

Do not treat semantic agreement as the percentage of clinically correct notes, or its complement as the error rate. The metric's algorithm, comparison reference, analytic denominator and weighting method remain unverified. A similarly named vendor metric may not measure the same thing.

02 Scale is not the audit denominator

The consultation population and detailed quality sample are different. Do not present the millions of consultations as millions of manually audited notes. Readability and structured-quality results do not establish patient outcomes or causal clinical benefit.

03 Outpatient evidence is not a hospital-wide license

US language, specialty mix, EHR workflows, documentation requirements and reimbursement arrangements require local evaluation. Inpatient and emergency workflows are separate expansion decisions, not automatic extensions of an outpatient test.

04 Capacity is not cash

The abstract supplies no measured ROI. Clinician time multiplied by salary is not cash savings unless spending actually falls. Schedules, rooms, staffing and demand must allow recovered time to become completed visits; collected contribution and paid expense changes must be verified.

05 Workforce relief is its own benefit

Reduced evening documentation can justify investment as workforce relief. It does not require a fictional revenue claim. Authorization for workforce relief and authorization for financial return should state different evidence and objectives.

Interpretation boundary: The cautions and investment recommendations above are ALHQ editorial interpretation. They are not attributed to an unread study limitations section. No unverified BMJ or preprint findings are used as supporting evidence.



From Evidence to a Local Decision

The scorecard is an ALHQ proposal, not a validated instrument. Its seven measures cover usable adoption, clinical content, net documentation work, realized capacity, realized financial benefit, incremental cost and control exceptions. Local owners set criteria before seeing results.

Illustrative 30-day evaluation

Days 1-5 | Fix the comparison

Select two outpatient specialties. Establish baseline with the same clinicians' prior encounters and, where feasible, a concurrent non-scribe group. Record case mix, language, staffing and schedule changes. Define severity, escalation, mandatory review, costs and break-even assumptions.

Days 6-25 | Test drafts and released notes

Maintain clinician review of every draft before release. The proposed starting audit is 100 scribe-assisted and 100 comparator notes, stratified across clinicians, specialties, languages and complexity. Report oversampled risky encounters separately; double-review a subset. This is not a statistically validated sample or rare-event safety certification.

Days 26-30 | Reconcile, then decide

Compare quality, work, adoption, cost and capacity. Separate early training effects and comparison-group differences from causal attribution. Collections may remain immature: mark finance provisional and set a reconciliation date. A short operational evaluation is not a patient-outcome trial.

Before broadening the evidence claim

Obtain and read the full methods and limitations, confirm the metric definition and audit procedures, and verify financial or patient endpoints directly before asserting them. Do not replace the supplied journal version with an earlier preprint. Retrieve failures do not establish that a publication does not exist.

Decision: Expand only in tested workflows that meet pre-agreed local criteria and have an explicit, affordable benefit case. Continue within the cap to close defined gaps. Pause and escalate on local serious-error/privacy triggers, failed mandatory review or invalidating version changes, regardless of aggregate rates.